



Exercise 5.3: Using ConfigMaps Configure Ambassador Containers

In an earlier lab we added a second Ambassador container to handle logging. Now that we have learned about using ConfigMaps and attaching storage we will use configure our basic pod.

1. Review the YAML for our earlier simple pod. Recall that we added an Ambassador style logging container to the pod but had not fully configured the logging.

```
student@ckad-1:~$ cat basic.yaml
```

```
<output_omitted>
containers:
- name: webcont
  image: nginx
  ports:
  - containerPort: 80
- name: fdlogger
  image: fluent/fluentd
```

2. Let us begin by adding shared storage to each container. We will use the `hostPath` storage class to provide the PV and PVC. First we create the directory.

```
student@ckad-1:~$ sudo mkdir /tmp/weblog
```

3. Now we create a new PV to use that directory for the `hostPath` storage class. We will use the `storageClassName` of `manual` so that only PVCs which use that name will bind the resource.

```
student@ckad-1:~$ vim weblog-pv.yaml
```

YAML

weblog-pv.yaml

```
1 kind: PersistentVolume
2 apiVersion: v1
3 metadata:
4   name: weblog-pv-volume
5   labels:
6     type: local
7 spec:
8   storageClassName: manual
9   capacity:
10    storage: 100Mi
11  accessModes:
12    - ReadWriteOnce
13  hostPath:
14    path: "/tmp/weblog"
```

4. Create and verify the new PV exists.

```
student@ckad-1:~$ kubectl create -f weblog-pv.yaml
```

```
persistentvolume/weblog-pv-volume created
```

```
student@ckad-1:~$ kubectl get pv weblog-pv-volume
```

NAME	CAPACITY	ACCESS MODES	RECLAIM POLICY
STATUS	CLAIM	STORAGECLASS	REASON AGE
weblog-pv-volume	100Mi	RWO	Retain
Available		manual	21s

5. Next we will create a PVC to use the PV we just created.

```
student@ckad-1:~$ vim weblog-pvc.yaml
```

YA
ML

weblog-pvc.yaml

```
1 kind: PersistentVolumeClaim
2 apiVersion: v1
3 metadata:
4   name: weblog-pv-claim
5 spec:
6   storageClassName: manual
7   accessModes:
8     - ReadWriteOnce
9   resources:
10    requests:
11      storage: 100Mi
```

6. Create the PVC and verify it shows as Bound to the the PV we previously created.

```
student@ckad-1:~$ kubectl create -f weblog-pvc.yaml
```

```
persistentvolumeclaim/weblog-pv-claim created
```

```
student@ckad-1:~$ kubectl get pvc weblog-pv-claim
```

NAME	STATUS	VOLUME	CAPACITY	ACCESS MODES
STORAGECLASS	AGE			
weblog-pv-claim	Bound	weblog-pv-volume	100Mi	RWO
manual	79s			

7. We are ready to add the storage to our pod. We will edit three sections. The first will declare the storage to the pod in general, then two more sections which tell each container where to make the volume available.

```
student@ckad-1:~$ vim basic.yaml
```

YA
ML

basic.yaml

```
1 apiVersion: v1
2 kind: Pod
3 metadata:
4   name: basicpod
5   labels:
6     type: webserver
7 spec:
8   volumes:                                     #<-- Add three lines, same depth as containers
9     - name: weblog-pv-storage
10       persistentVolumeClaim:
11         claimName: weblog-pv-claim
12   containers:
13     - name: webcont
14       image: nginx
15       ports:
16     - containerPort: 80
```



```

17     volumeMounts:                                #<-- Add three lines, same depth as ports
18       - mountPath: "/var/log/nginx/"
19         name: weblog-pv-storage                    # Must match volume name above
20   - name: fdlogger
21     image: fluent/fluentd
22     volumeMounts:                                #<-- Add three lines, same depth as image:
23       - mountPath: "/var/log"
24         name: weblog-pv-storage                    # Must match volume name above

```

8. At this point we can create the pod again. When we create a shell we will find that the `access.log` for `nginx` is no longer a symbolic link pointing to `stdout` it is a writable, zero length file. Leave a **tail** of the log file running.

```
student@ckad-1:~$ kubectl create -f basic.yaml
```

```
pod/basicpod created
```

```
student@ckad-1:~$ kubectl exec -c webcont -it basicpod -- /bin/bash
```



On Container

```

root@basicpod:/# ls -l /var/log/nginx/access.log
-rw-r--r-- 1 root root 0 Oct 18 16:12 /var/log/nginx/access.log

root@basicpod:/# tail -f /var/log/nginx/access.log

```

9. Open a second connection to your node. We will use the pod IP as we have not yet configured a service to expose the pod.

```
student@ckad-1:~$ kubectl get pods -o wide
```

```

NAME          READY STATUS   RESTARTS  AGE    IP             NODE
NOMINATED NODE
basicpod 2/2   Running    0          3m26s  192.168.213.181 ckad-1
<none>

```

10. Use **curl** to view the welcome page of the webserver. When the command completes you should see a new entry added to the log. Right after the GET we see a 200 response indicating success. You can use **ctrl-c** and **exit** to return to the host shell prompt.

```

student@ckad-1:~$ curl http://192.168.213.181
<!DOCTYPE html>
<html>
<head>
<title>Welcome to nginx!</title>
<output_omitted>

```



On Container

```
192.168.32.128 - - [18/Oct/2018:16:16:21 +0000] "GET / HTTP/1.1" 200 612 "-" "curl/7.47.0" "-"
```

11. Now that we know the `webcont` container is writing to the PV we will configure the logger to use that directory as a source. For greater flexibility we will configure **fluentd** using a `configMap`. The details of the data settings can be found in **fluentd** documentation here: <https://docs.fluentd.org/v1.0/categories/config-file>

```
student@ckad-1:~$ vim weblog-configmap.yaml
```



weblog-configmap.yaml

```

1  apiVersion: v1
2  kind: ConfigMap
3  metadata:
4    name: fluentd-config
5  data:
6    fluentd.conf: |
7      <source>
8        @type tail
9        format none
10       path /var/log/nginx/access.log
11       tag count.format1
12     </source>
13
14     <match *.*>
15       @type forward
16
17       <server>
18         name localhost
19         host 127.0.0.1
20       </server>
21     </match>

```

12. Create the new configMap.

```

student@ckad-1:~$ kubectl create -f weblog-configmap.yaml
configmap/fluentd-config created

```

13. Now we will edit the pod yaml file so that the **fluentd** container will mount the configmap as a volume and reference the variables inside the config file. You will add three areas, the volume declaration to the pod, the env parameter and the mounting of the volume to the fluentd container

```

student@ckad-1:~$ vim basic.yaml

```



basic.yaml

```

1  ....
2  volumes:
3    - name: weblog-pv-storage
4      persistentVolumeClaim:
5        claimName: weblog-pv-claim
6    - name: log-config                                #<-- This and two lines following
7      configMap:
8        name: fluentd-config                          # Must match existing configMap
9  ....
10  image: fluent/fluentd
11  env:                                                #<-- This and two lines following
12    - name: FLUENTD_ARGS
13      value: -c /etc/fluentd-config/fluentd.conf
14  ....
15  volumeMounts:
16    - mountPath: "/var/log"
17      name: weblog-pv-storage
18    - name: log-config                                #<-- This and next line
19      mountPath: "/etc/fluentd-config"

```

14. At this point we can delete and re-create the pod. If we had a listening agent running on **localhost**, where the we messages are forwarded as declared in the configMap, we would see access messages.

```
student@ckad-1:~$ kubectl delete pod basicpod
pod "basicpod" deleted
```

```
student@ckad-1:~$ kubectl create -f basic.yaml
pod/basicpod created
```

```
student@ckad-1:~$ kubectl get pod basicpod
```

NAME	READY	STATUS	RESTARTS	AGE
basicpod	2/2	Running	0	8s

15. Look at the logs for both containers. You should see some output for the fdlogger but not for webcont.

```
student@ckad-1:~$ kubectl logs basicpod webcont
```

```
student@ckad-1:~$ kubectl logs basicpod fdlogger
```

```
2019-03-06 18:55:33 +0000 [info]: parsing config file is succeeded path="/fluentd/etc/fluent.conf"
2019-03-06 18:55:33 +0000 [warn]: [output_docker1] 'time_format' specified without 'time_key', will be ignored
2019-03-06 18:55:33 +0000 [warn]: [output1] 'time_format' specified without 'time_key', will be ignored
2019-03-06 18:55:33 +0000 [info]: using configuration file: <ROOT>
  <source>
    @type forward
    @id input1
  <output_omitted>
```